

Q2. Consider natural-joining tables R(a, b) and S(a,c). Suppose we have the following scenario.

- i. R is a clustered relation with 10,000 blocks.
- ii. S is a clustered relation with 20,000 blocks.
- iii. 102 pages available in main memory for the join.
- iv. Assume the output of join is given to the next operator in the query execution plan (instead of writing to the disk) and thus the cost of writing the output is ignored.

Describe the steps for each of the following join algorithms.

For sorting and hashing-based algorithms, also indicate the sizes of output from each step.

What is the total number of block I/O's needed for each algorithm?

Which algorithm is most efficient in terms of block's I/O?

$$M = 102$$

$$M-2 = 100$$

$$B(R) = 10000 \text{ blocks}$$

$$B(S) = 20000 \text{ blocks}$$

a. [10 points] (Block-based) nested-loop join with R as the outer relation.

Solution

A: $R \bowtie S$

Pseudocode:

```
For each (102-2) blocks br of R do:
  For each block bs of S do:
    For each tuple r in br do:
      For each tuple s in bs do:
        If "r and s join" then output (r, s)
```

Number of block I/O's:

- 1. Read R once, cost $B(R) = 10000$ blocks
- 2. In main memory it can load 100 blocks at a time when outer loop is executed.
- 3. R as the outer relation, so it will have $B(R)/(M-2) = 10000/100 = 100$ time iterations in outer loop.

Each time, 100 blocks are loaded into main memory

- 4. In each iteration, make one pass through R, scanning entire S to check every block in S and join tuples to output buffer.

So the cost is $(B(R)/(M-2)) * B(S) = 100 * 20000 = 2000000$

Total Cost: $B(R) + ((B(R)/(M-2)) * B(S)) = 10000 + ((10000/100) * 20000) = \mathbf{2010000 \text{ block I/O's}}$

b. [10 points] (Block-based) nested-loop join with S as the outer relation.

Solution

$$M = 102$$

$$M-2 = 100$$

$B(R) = 10000$ blocks

$B(S) = 20000$ blocks

Pseudocode:

For each (102-2) blocks bs of S do:

 For each block br of R do:

 For each tuple s in bs do:

 For each tuple r in br do:

 If " s and r join" then output (s, r)

Number of block I/O's:

1. S as the outer relation, make a pass through S , cost $B(S) = 20000$ blocks.

2. In outer loop, it will execute $B(S)/(M-2) = 20000/100 = 200$ times. Each time, 200 blocks are loaded into main memory.

3. In inner loop, each time scanning R in input buffer to check every block in R and join tuples to output buffer.

The cost is $(B(S)/(M-2))*B(R) = (20000/100)*10000 = 200*10000 = 2000000$

4. Total Cost: $B(S) + ((B(S)/(M-2))*B(R)) = 20000 + ((20000/100)*10000) = \mathbf{2020000}$ block I/O's

c. [20 points] Sort-merge join (assume only 100 pages are used for sorting and 101 pages for merging). Note that if join cannot be done by using only a single merging pass, runs from one or both relations need to be further merged, in order to reduce the number of runs. Select the relation with a larger number of runs for further merging first if both have too many runs.

Solution

$M=101$ blocks

$B(R) = 10000$ blocks

$B(S) = 20000$ blocks

-> 100 pages are used for sorting

so there are:

$B(R)/100 = 10000/100 = 100$ runs (each size 100) of R

$B(S)/100 = 20000/100 = 200$ runs (each size 100) of S

The cost (sort and write) is $2B(R)+2B(S) = 2*10000 + 2*20000 = 20000 + 40000 = 60000$ blocks.

Send them back to disk.

-> merging - 1 = 101-1 = 100,

while $B(R)+B(S) = 10000 + 20000 = 30000$

results into $B(R)+B(S) \geq (M-1)M \implies 30000 \geq (101-1)*100 \implies 30000 \geq 10000$.

So, we cannot merge $M-1$ runs from R and S directly

We need to merge:

R into $100/100 = 1$ run, Cost = $4B(R)$

S into $200/100 = 2$ runs Cost = $4B(S)$

-> Join by merging 1 run with 2 runs, Cost = $B(R) + B(S)$

-> Total = $5B(R) + 5B(S) = 5(B(R) + B(S)) = 5 \cdot (30000) = \mathbf{150000}$

d. [20 points] Partitioned-hash join (assume 101 pages used in partitioning of relations and no hash table is used to lookup in joining tuples). Note if buckets are still too large to join within memory, you should further partition them.

Solution

$B(R) = 10000$ blocks

$B(S) = 20000$ blocks

-->

Hash R into $M-1(100)$ buckets, 100 blocks/bucket(R_1)

Hash S into $M-1(100)$ buckets, 200 blocks/bucket(S_1)

Cost : $2 \cdot B(R) + 2 \cdot B(S)$

--> Extra hash is required as sum of bucket R and S is more than 100

Hash R into 100 buckets, 1 block/bucket(R_{11})

Hash S into 100 buckets, 2 blocks/bucket(S_{11})

Cost : $2 \cdot B(R) + 2 \cdot B(S)$

-> Join by merging R with S, Cost = $B(R) + B(S)$

-> Total = $5B(R) + 5B(S) = 5(B(R) + B(S)) = 5 \cdot (30000) = \mathbf{150000}$

Q. Which algorithm is most efficient in terms of block's I/O?

The Sort-merge join and Partitioned-hash join both are the most efficient in terms of block's I/O.